

DATA621_HW1 Spring'18

Jaan Bernberg

February 9, 2018

Purpose

This assignment explores, analyzes, and models data from baseball team seasons from 1871-2006 inclusive (stats adjusted to match 162-game season). Each record represents the performance of a given baseball team per year.

The goal is to create a model that can predict baseball team wins based on the performance metrics captured in the training dataset `moneyball-training-data.csv` and then review the predictive quality by running the model with `moneyball-evaluation-data.csv` the testing dataset.

Data Exploration

High-level Findings

- There are ~2300 observations with 16 variables included
- Five of the variables have a number of missing values
- Two of the variables have outliers
- Multicollinearity present in seven variables
- Untransformed (raw) data is not highly correlated with response

Findings Detail

The below investigation was completed prior to model-fitting in order to get a sense for the density and quality of the training data set and to preview what types of actions may need to be taken in the model-fitting steps.

Summary Statistics:

Table 1: Selected Stats

	vars	n	mean	sd	median	min	max	range	kurtosis	IQR
TARGET_WINS	1	2276	80.79	15.75	82.0	0	146	146	1.03	21.00
TEAM_BATTING_H	2	2276	1469.27	144.59	1454.0	891	2554	1663	7.28	154.25
TEAM_BATTING_2B	3	2276	241.25	46.80	238.0	69	458	389	0.01	65.00
TEAM_BATTING_3B	4	2276	55.25	27.94	47.0	0	223	223	1.50	38.00
TEAM_BATTING_HR	5	2276	99.61	60.55	102.0	0	264	264	-0.96	105.00
TEAM_BATTING_BB	6	2276	501.56	122.67	512.0	0	878	878	2.18	129.00
TEAM_BATTING_SO	7	2174	735.61	248.53	750.0	0	1399	1399	-0.32	382.00
TEAM_BASERUN_SB	8	2145	124.76	87.79	101.0	0	697	697	5.49	90.00
TEAM_BASERUN_CS	9	1504	52.80	22.96	49.0	0	201	201	7.62	24.00
TEAM_BATTING_HBP	10	191	59.36	12.97	58.0	29	95	66	-0.11	16.50
TEAM_PITCHING_H	11	2276	1779.21	1406.84	1518.0	1137	30132	28995	141.84	263.50

	vars	n	mean	sd	median	min	max	range	kurtosis	IQR
TEAM_PITCHING_HR	12	2276	105.70	61.30	107.0	0	343	343	-0.60	100.00
TEAM_PITCHING_BB	13	2276	553.01	166.36	536.5	0	3645	3645	96.97	135.00
TEAM_PITCHING_SO	14	2174	817.73	553.09	813.5	0	19278	19278	671.19	353.00
TEAM_FIELDING_E	15	2276	246.48	227.77	159.0	65	1898	1833	10.97	122.25
TEAM_FIELDING_DP	16	1990	146.39	26.23	149.0	52	228	176	0.18	33.00

Missing Values:

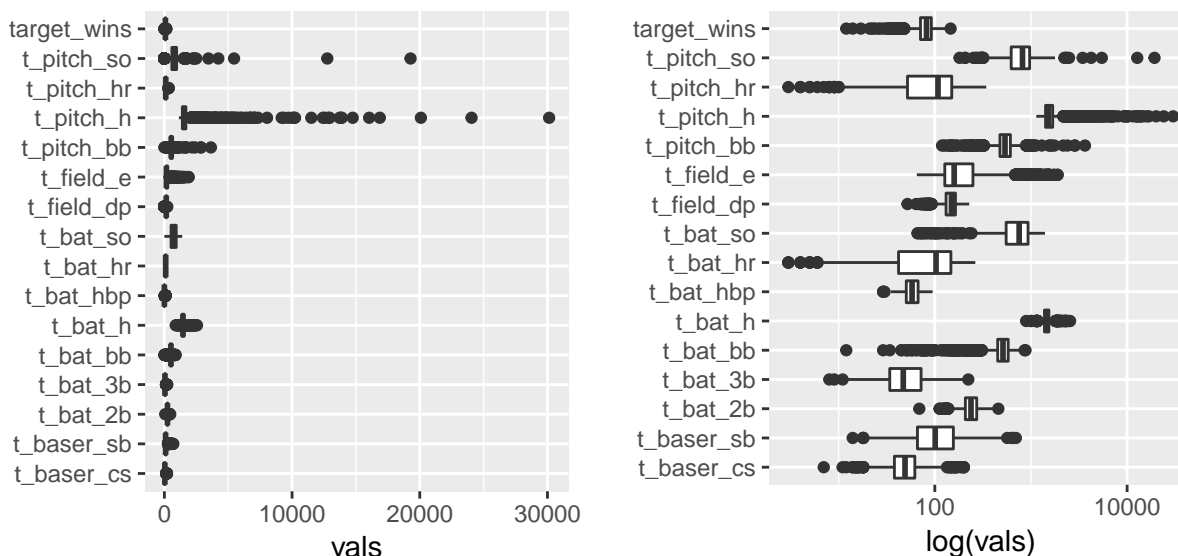
Review of the NA's (see column n, Table1) revealed the following sparse variables. The value after the variable name represents the number of missing values followed % density on 2,276 records available in most (11 of 16) variables:

- TEAM_BATTING_HBP: 2085: 8%
- TEAM_BASERUN_CS: 772: 66%
- TEAM_FIELDING_DP: 286: 87%
- TEAM_PITCHING_SO: 102: 96%
- TEAM_BATTING_SO: 102: 96%

Outliers:

Below, a preview of each variable in the data using box plots reveals at least 2 variables containing points with significant outliers:


- TEAM_PITCHING_H (hits allowed by pitchers)
- TEAM_PITCHING_SO (strikeouts by pitchers)



Note that the right-hand side plot is a duplicate of the left-hand plot but uses a log scale for the value-axis.

Multicollinearity

The below heat-map style plot highlights the correlation that exists between some of the predictor variables in the training set:



For convenience, the above plot was built by substituting 0's for NA's for all data but model fits do not. The primary takeaway is that in isolation, none of the variables appears to be a great predictor of `TARGET_WINS`.

Data Preperation

Several different methods were investigated. The below provides a summary of methods used to prepare each of the data-sets for 3 different model fits for the purpose of prediction.

Below, I have outlined the steps taken to prepare individual data-sets for each model.

Model 1 Data:

Motivation: Based on the most complete set of records with minimal imputation:

- For `TEAM_BATTING_SO` and `TEAM_PITCHING_SO`, NA-values replaced with the `mean` and `median` respectively to address normal and skewed distributions respectively - these variables were only missing 4% of their records.
- Removed `TEAM_BATTING_HBP` because of its numerous missing values (92% of records missing values for this variable.)

Model 2 Data:

Motivation: a refinement of the first model data-set that excludes variables that may be doing 'double-duty':

- Replace NA with `mean` for: `TEAM_FIELD_DP` to cover 13% of missing values
- Replace NA with `median` for `TEAM_PITCH_SO` to cover 4% missing values

- Remove TEAM_BATTING_SO because it's highly correlated with TEAM_PITCHING_SO
- Remove TEAM_BASERUN_CS and TEAM_BATTING_HBP because 44% and 92% of records are missing respectively.

Model3 Data:

Motivation: Address all/most multicollinearity encountered and reduce the number of predictors overall to simplify the model. Variable selection between correlated predictors was based on 'hitting-variables' because those intuitively may result in more runs, and therefore, more wins.

- Remove TEAM_BASERUN_CS and TEAM_BATTING_HBP because 44% and 92% of records are missing respectively.
 - Remove TEAM_BATTING_2B/3B/HR as TEAM_BATTING_H covers all hit types - this is an attempt to simplify the model fit
 - Remove one side of the correlated pairs focusing on batting-related variables such that:
 - TEAM_BATTING_SO kept in, removed TEAM_PITCHING_SO
 - TEAM_BATTING_H kept in, removed TEAM_PITCHING_H
 - TEAM_PITCHING_HR kept in, removed TEAM_BATTING_HR as TEAM_BATTING_H covers all hits
 - TEAM_FIELDING_E kept in, removed TEAM_FIELDING_DP
-

Build Models

Model Fitting

Models were fit with the respective data-sets created in the previous section. Each section will display the resulting fit's coefficients but the discussion of the coefficients will be saved for the *Model Selection* section.

Fit1:

Fit1 utilized the `stepAIC` function from the MASS package to perform step-wise model selection by AIC. I used this auto-selection method to investigate an off-the-shelf model's performance with no additional transformations or imputation beyond what was originally discussed.

Table 3: Model 1 coefficients

Vars	Coefs.
(Intercept)	58.4461
t_bat_h	0.0255
t_bat_2b	-0.0698
t_bat_3b	0.1616
t_bat_hr	0.0977
t_bat_bb	0.0395
t_baser_sb	0.0360
t_baser_cs	0.0518
t_pitch_h	0.0091
t_pitch_so	-0.0208
t_field_e	-0.1560
t_field_dp	-0.1131

Model 2 Fit:

An attempt at employing the `powerTransform` function from the `car` package to see if box-cox would yeild any transformation but my result was $\lambda \approx 1$ so I did not transform the model.

Table 4: Model 2 coefficients

Vars	Coefs.
(Intercept)	29.7307
t_bat_h	0.0422
t_bat_2b	-0.0490
t_bat_3b	0.0741
t_bat_hr	0.0147
t_bat_bb	0.0303
t_baser_sb	0.0565
t_pitch_h	0.0027
t_pitch_hr	0.0557
t_pitch_bb	-0.0030
t_pitch_so	-0.0102
t_field_e	-0.0517
t_field_dp	-0.1168

Model 3 Fit:

The result of the box-cox transform here was a value of $\lambda \approx 1$ which again called for no transformation.

Table 5: Model 3 coefficients

Vars	Coefs.
(Intercept)	24.3922
t_bat_h	0.0343
t_bat_bb	0.0239
t_bat_so	-0.0160
t_baser_sb	0.0558
t_pitch_hr	0.0681
t_pitch_bb	-0.0034
t_field_e	-0.0346

Select Models

Criteria:

My selection criteria is based on comparing R^2 and F-Statistics. If the performance is similar across these measures, I will select the model that is the easiest to describe of the two.

Model Selection:

I utilized the `stargazer` package to create the table below that compares Models 1, 2, and 3. Note that at the bottom of the table the models are compared by R^2 , adjusted R^2 , and the F -statistic.

```
##
## =====
##                               Dependent variable:
## -----
##                               target_wins
##                               (1)         (2)         (3)
## -----
## t_bat_h          0.026*** (0.006)      0.042*** (0.004)      0.034*** (0.003)
## t_bat_2b        -0.070*** (0.009)     -0.049*** (0.009)
## t_bat_3b         0.162*** (0.022)      0.074*** (0.017)
## t_bat_hr         0.098*** (0.009)       0.015 (0.027)
## t_bat_bb         0.039*** (0.003)      0.030*** (0.006)      0.024*** (0.004)
## t_bat_so         -0.016*** (0.002)
## t_baser_sb       0.036*** (0.009)      0.057*** (0.004)      0.056*** (0.004)
## t_baser_cs       0.052*** (0.018)
## t_pitch_h        0.009*** (0.002)      0.003*** (0.0004)
## t_pitch_hr       0.056** (0.025)       0.068*** (0.008)
## t_pitch_bb       -0.003 (0.004)        -0.003 (0.003)
## t_pitch_so      -0.021*** (0.002)     -0.010*** (0.002)
## t_field_e        -0.156*** (0.010)     -0.052*** (0.003)      -0.035*** (0.003)
## t_field_dp       -0.113*** (0.013)     -0.117*** (0.012)
## Constant         58.446*** (6.589)     29.731*** (5.187)      24.392*** (4.803)
## -----
## Observations      1,486                2,145                2,043
## R2                0.438                0.374                0.337
## Adjusted R2       0.434                0.370                0.335
## F Statistic    104.596*** (df = 11; 1474) 106.144*** (df = 12; 2132) 147.644*** (df = 7; 2035)
## =====
## Note:                                                     *p<0.1; **p<0.05; ***p<0.01
```

Observations from these 3 models:

- None of the 3 models have very high R^2 values.
- Each model's F-statistic indicates that there is a relationship between predictor and the response variables included in the fits

Coefficients:

Model1: The intercept ('Constant') is too low for the mean value of `TARGET_WINS` but it's not likely that every baseball team will get ~58 wins a season. The signs of the coefficients are problematic for several cases e.g. `TEAM_BASERUN_CS` (caught stealing) is positive but that would be detrimental to the wins and `TEAM_FIELDING_E` (fielding errors) are negative but fielding errors would lead to more on-bases and therefore, more runs.

Model2: More-intuitive intercept value but similarly un-intuitive coefficient signs that mirror the issues pointed out in Model1.

Model3: This model contains the fewest variables but 2 of them still have problematic signs including `TEAM_FIELDING_E` and `TEAM_PITCHING_BB` which I would expect to be positively correlated with wins.

Based on my criteria, since the performance of Model 1 is superior compared to the three models, I have selected **Model 1**.

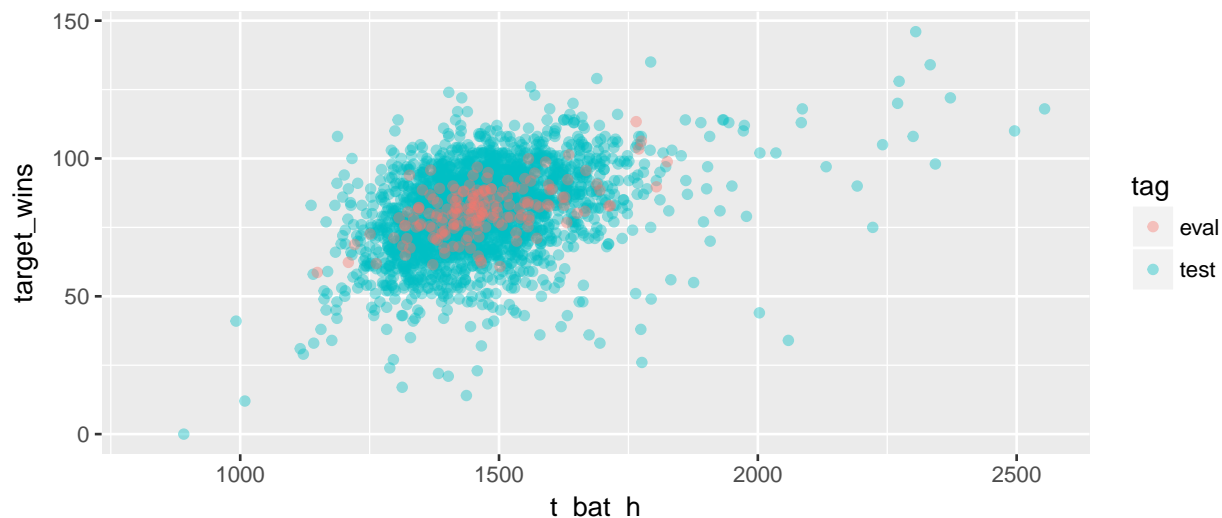
Model Evaluation:

```
##
## Call:
## lm(formula = target_wins ~ t_bat_h + t_bat_2b + t_bat_3b + t_bat_hr +
##      t_bat_bb + t_baser_sb + t_baser_cs + t_pitch_h + t_pitch_so +
##      t_field_e + t_field_dp, data = f1_bsb_raw[, -1])
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -30.5830  -6.7313  -0.1643   6.5323  27.8502
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  58.446058   6.588864   8.870 < 2e-16 ***
## t_bat_h       0.025502   0.005837   4.369 1.34e-05 ***
## t_bat_2b     -0.069828   0.009292  -7.515 9.86e-14 ***
## t_bat_3b      0.161620   0.022159   7.294 4.91e-13 ***
## t_bat_hr      0.097746   0.009435  10.360 < 2e-16 ***
## t_bat_bb      0.039484   0.003356  11.765 < 2e-16 ***
## t_baser_sb    0.036034   0.008668   4.157 3.41e-05 ***
## t_baser_cs    0.051768   0.018200   2.844 0.004511 **
## t_pitch_h     0.009070   0.002353   3.855 0.000121 ***
## t_pitch_so   -0.020827   0.002312  -9.006 < 2e-16 ***
## t_field_e    -0.155970   0.009917 -15.728 < 2e-16 ***
## t_field_dp   -0.113149   0.013121  -8.623 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 9.548 on 1474 degrees of freedom
## (790 observations deleted due to missingness)
## Multiple R-squared:  0.4384, Adjusted R-squared:  0.4342
## F-statistic: 104.6 on 11 and 1474 DF,  p-value: < 2.2e-16
```

The summary output from Model 1, above, reveals a number of challenges:

- The **Intercept** value of ~ 60 could be too high given that **TARGET_WINS** can take values as low as 0 (team loses all games) but it could also be too low given that the **mean** and **median** is *approx* 80.
- The coefficients have very low p-values so the observed data is extremely unlikely to have arisen if the null hypothesis were true (i.e. that $H_0 : \beta_i = 0$ so we can reject the null hypothesis)
- 790 observations related to the missing values in columns **TEAM_BATTING_HBP**
- The R^2 and multiple R^2 values are low ($R^2 < .70$) at $\approx .43$ - this indicates to me that the data isn't very close to the fitted regression and may provide poor predictions
- The negative coefficients under **TEAM_BATTING_2B** and other base-gaining activities is counter-intuitive as getting on-base intuitively would result in more wins.

Below, the residual and QQ plots of *Model1* appear to support a normal classification with some reservations. The residuals vs. the fitted plot appear to have a slight 'kink' when I expected an even, straight line. The normal QQ plot doesn't deviate from the diagonal except in a few outlying cases.



R Code for Assignment 1:

```
#confint(fitA)

knitr::opts_chunk$set(echo = TRUE)
#library(memisc)
library(stargazer)
library(tidyverse)
#library(GGally)
# read in data.
bsb_train_raw <- read.csv("moneyball-training-data.csv")
mini_tbl<-psych::describe(bsb_train_raw[, -1], IQR=T)[,c(1:5,8:10,12,14)]

bsb_train <- data.table::setnames(bsb_train_raw,
                                tolower(names(bsb_train_raw[1:17])))
bsb_train[is.na(bsb_train)] <- 0 # used for simple plots - not used for model fitting
                                # see data preperation section for more details.

col_name_updater <- function(x, text_to_updt, replc_text){
  colnames(x) = gsub(text_to_updt, replc_text, colnames(x))
  return(x)
}
# shorten and clean colnames
c1 <- col_name_updater(bsb_train, "batting", "bat")
c2 <- col_name_updater(c1, "pitching", "pitch")
c3 <- col_name_updater(c2, "fielding", "field")
c4 <- col_name_updater(c3, "baserun", "baser")
bsb_data <- col_name_updater(c4, "team", "t")
rm(c1,c2,c3, c4)

c1 <- col_name_updater(bsb_train_raw, "batting", "bat")
c2 <- col_name_updater(c1, "pitching", "pitch")
c3 <- col_name_updater(c2, "fielding", "field")
c4 <- col_name_updater(c3, "baserun", "baser")
bsb_train_raw <- col_name_updater(c4, "team", "t")
```

```

rm(c1,c2,c3, c4)
knitr::kable(round(mini_tbl,2), caption = "Selected Stats")
rm(mini_tbl)

p1 <- bsb_data %>%
  gather(-index, key = "vars", value = "vals") %>%
  select_("vars", "vals") %>%
  ggplot() +
  geom_boxplot(aes(x = vars, y=vals)) +
  coord_flip() + xlab("")

p2 <- bsb_data %>%
  gather(-index, key = "vars", value = "vals") %>%
  select_("vars", "vals") %>%
  ggplot() +
  geom_boxplot(aes(x = vars, y=vals)) +
  scale_y_log10() +
  coord_flip() + ylab("log(vals)") + xlab("")

suppressMessages(gridExtra::grid.arrange(p1, p2, nrow=1))

# check correlations between prediction vars:
GGally::ggcorr(bsb_data[,c(6,8,12,13,15:17)],
  geom = "circle", nbreaks = 6,
  label=T, label_size = 3)

# see if there's cor between target wins and others...
sm_tab <- cor(bsb_data[, -1]) %>% as.data.frame() %>%
  tibble::rownames_to_column(var = "var_name") %>%
  mutate(cor_val = round(target_wins , 4),
    abs_cor_val = abs(cor_val)) %>%
  select_("var_name", "abs_cor_val", "cor_val") %>%
  arrange(desc(abs_cor_val)) %>%
  na.omit()

knitr::kable(head(sm_tab), caption = "Abs. Correlation Between Target Wins and other variables")
bsb_data[, -1] %>%
  gather(-target_wins, key = "vars", value = "vals") %>%
  ggplot(aes(x = vars, y = target_wins)) +
  geom_point(alpha=.1) +
  facet_wrap(~ vars, scales = "free", ncol = 3)
na_ifelse <- function(x, metric){
  #function to replace NAs
  if(metric=="mean"){
    result = ifelse(is.na(x), mean(x, na.rm=T), x)
  }else if (metric == "med"){
    result = ifelse(is.na(x), median(x, na.rm=T), x)
  }else if (metric == "min"){
    result = ifelse(is.na(x), min(x, na.rm=T), x)
  }else
    result = ifelse(is.na(x), 0, x)
  return(result)
}

f1_bsb_raw <- bsb_train_raw %>%
  mutate(t_bat_so = na_ifelse(t_bat_so, "mean"),

```

```

    t_pitch_so = na_ifelse(t_pitch_so, "med")) %>%
  select(-t_bat_hbp)
# sapply(new_bsb_raw, function(x) sum(is.na(x))) # count na's per column!
# View(new_bsb_raw)
f2_bsb_raw <- bsb_train_raw %>%
  mutate(t_field_dp = na_ifelse(t_field_dp, "mean"),
         t_pitch_so = na_ifelse(t_pitch_so, "med")) %>%
  select(-t_bat_so, -t_baser_cs, -t_bat_hbp)
f3_bsb_raw <- bsb_train_raw %>%
  select(-t_baser_cs, -t_bat_hbp,
        -t_bat_2b, -t_bat_3b, -t_bat_hr,
        -t_pitch_so, -t_pitch_h, -t_field_dp)

## A function to pull the coefficients from a linear model and present them nicely.
get_coefs <- function(an_lm){
  temp = an_lm$coefficients %>% as.data.frame()
  temp$Vars = rownames(temp)
  colnames(temp)[1] = "Coefs."
  rownames(temp) = NULL
  return(temp[,2:1])
}

##### MODEL 1 FIT #####
fitA <- lm(target_wins ~., f1_bsb_raw[,-1])
fitA_stepped <- MASS::stepAIC(fitA, trace=F)

knitr::kable(get_coefs(fitA_stepped),
             caption = "Model 1 coefficients",
             digits = 4)
#summary(fitA_stepped)
par(mfrow=c(1,2))
plot(fitA_stepped, which=c(2,1))
par(mfrow=c(1,1))

##### MODEL 2 FIT #####
fitB <- lm(target_wins ~., f2_bsb_raw[,-1])
#fitB1 <- lm(target_wins+.01 ~., f2_bsb_raw[,-1])
#fitB_bc <- car::powerTransform(fitB1, family="bcPower") # 1.193639
#fitB_bc$roundlam # provides maximum estimator
#df <- cbind(lambda = fitB_bc$lambda, log_likli = fitB_bc$llik) %>% as.data.frame()
#df <- df %>% ungroup() %>% arrange(desc(log_likli))

#summary(fitB)
knitr::kable(get_coefs(fitB),
             caption = "Model 2 coefficients",
             digits = 4)
par(mfrow=c(1,2))
plot(fitB, which=c(2,1))
par(mfrow=c(1,1))

##### MODEL 3 FIT #####
fitC <- lm(target_wins ~., f3_bsb_raw[,-1])
#summary(fitC)
#fitC_bc <- lm(target_wins +1~., f3_bsb_raw[,-1])
#summary(fitC_bc)

```

```

#fitC1 <- lm(target_wins+1 ~., f3_bsb_raw[,-1])
#c1_bc <- car::powerTransform(fitC_bc, family="bcPower")
#c1_bc$roundlam # provides maximum estimator 1.275905
knitr::kable(get_coefs(fitC),
              caption = "Model 3 coefficients",
              digits = 4)
stargazer(fitA_stepped, fitB, fitC,
          type='text',
          column.sep.width = "1pt",
          single.row = T,
          omit.stat=c('ser'),
          no.space=TRUE)
summary(fitA_stepped)
par(mfrow=c(1,2))
plot(fitA_stepped, which=c(2,1))
par(mfrow=c(1,1))
confint(fitA_stepped)
#READ IN EVALUATION DATA
bsb_eval_raw <- read.csv("moneyball-evaluation-data.csv")
bsb_eval_raw <- data.table::setnames(bsb_eval_raw,
                                   tolower(names(bsb_eval_raw[1:16])))

c1 <- col_name_updater(bsb_eval_raw, "batting", "bat")
c2 <- col_name_updater(c1, "pitching", "pitch")
c3 <- col_name_updater(c2, "fielding", "field")
c4 <- col_name_updater(c3, "baserun", "baser")
bsb_eval_raw <- col_name_updater(c4, "team", "t")

bsb_eval_raw <- bsb_eval_raw %>%
  mutate(t_bat_so = na_ifelse(t_bat_so, "mean"),
         t_pitch_so = na_ifelse(t_pitch_so, "med")) %>%
  select(-t_bat_hbp)
# load test data into prediction
target_wins <- predict.lm(fitA_stepped, newdata=bsb_eval_raw[,-1])

# complete the test set
eval_set <- cbind(as.data.frame(target_wins), bsb_eval_raw[,-1]) %>%
  mutate(tag='eval')

tot_set <- f1_bsb_raw[,-1] %>%
  mutate(tag='test')

tot_set <- rbind(tot_set, eval_set)

x<-ggplot(tot_set, aes(x=t_bat_h, y=target_wins, color=tag)) +
  geom_point(alpha=.4)
x

# par(mfrow=c(1,2))
# plot(fitC)
# par(mfrow=c(1,1))
# confint(fitA)

```